

DATABASES

INTRODUCTION TO DATA SCIENCE

TIM KRASKA



DATA PROCESSING PIPELINES

Colin Mallows:

1. Identify data to collect and its relevance to your problem
2. Statistical specification of the problem
3. Method selection
4. Analysis of method
5. Interpret results for non-statisticians

Peter Huber:

1. Inspection
2. Error checking
3. Modification
4. Comparison
5. Modeling and model fitting
6. Simulation
7. What-if analyses
8. Interpretation
9. Presentation of conclusions

Ben Fry:

1. Acquire
2. Parse
3. Filter
4. Mine
5. Represent
6. Refine
7. Interact

Our Definition

1. Preparing to run a model
2. Running the model
3. Communicating the results

CLICKER QUESTION

How well do you know databases

- A. What are they?**
- B. I used a relational database in the past, but don't really know how they work.**
- C. I know SQL and tables**
- D. I know SQL, ER diagrams, and the relational algebra**
- E. I know normalization (e.g., 4th normal form) and, star and snowflake schemas**

WHY DATABASES

Why not store everything in flat files?

WHY DATABASES

Why not store everything in flat files?

- **Scalability → 100's of nodes**

WHY DATABASES

Why not store everything in flat files?

- Scalability → 100's of nodes
- Data redundancy and inconsistency

Name	Course	E-Mail	Grade
John Doe	CS112	jd@cs.brown.edu	B
C. Binnig	CS560		A
John Doe	CS560	John_doe@brown.edu	B

First Name	Last Name	Teaches	E-Mail	Grade
J.	Doe	CS112	jd@cs.brown.edu	B
Mike	Stonebraker	CS560	stonebraker@uni.edu	A
Carsten	Binnig		Carsten_binnig@brown.edu	B

Why is this a problem?

WHY DATABASES

Why not store everything in flat files?

- Scalability → 100's of nodes
- Data redundancy and inconsistency

Name	Course	E-Mail	Grade
John Doe	CS112	jd@cs.brown.edu	B
C. Binnig	CS560		A
John Doe	CS560	John_doe@brown.edu	B

First Name	Last Name	Teaches	E-Mail	Grade
J.	Doe	CS112	jd@cs.brown.edu	B
Mike	Stonebraker	CS560	stonebraker@uni.edu	A
Carsten	Binnig		Carsten_binnig@brown.edu	B

Why is this a problem?

- Wasted space (?)
- Potential inconsistencies

(e.g., multiple formats, John Smith vs Smith J.)

WHY DATABASES

Why not store everything in flat files?

- **Scalability** → 100's of nodes
- **Data redundancy and inconsistency**
- **Data retrieval**
 - Find the student who took CS18
 - Find the student who took CS18 and has a GPA > 3.5

WHY DATABASES

Why not store everything in flat files?

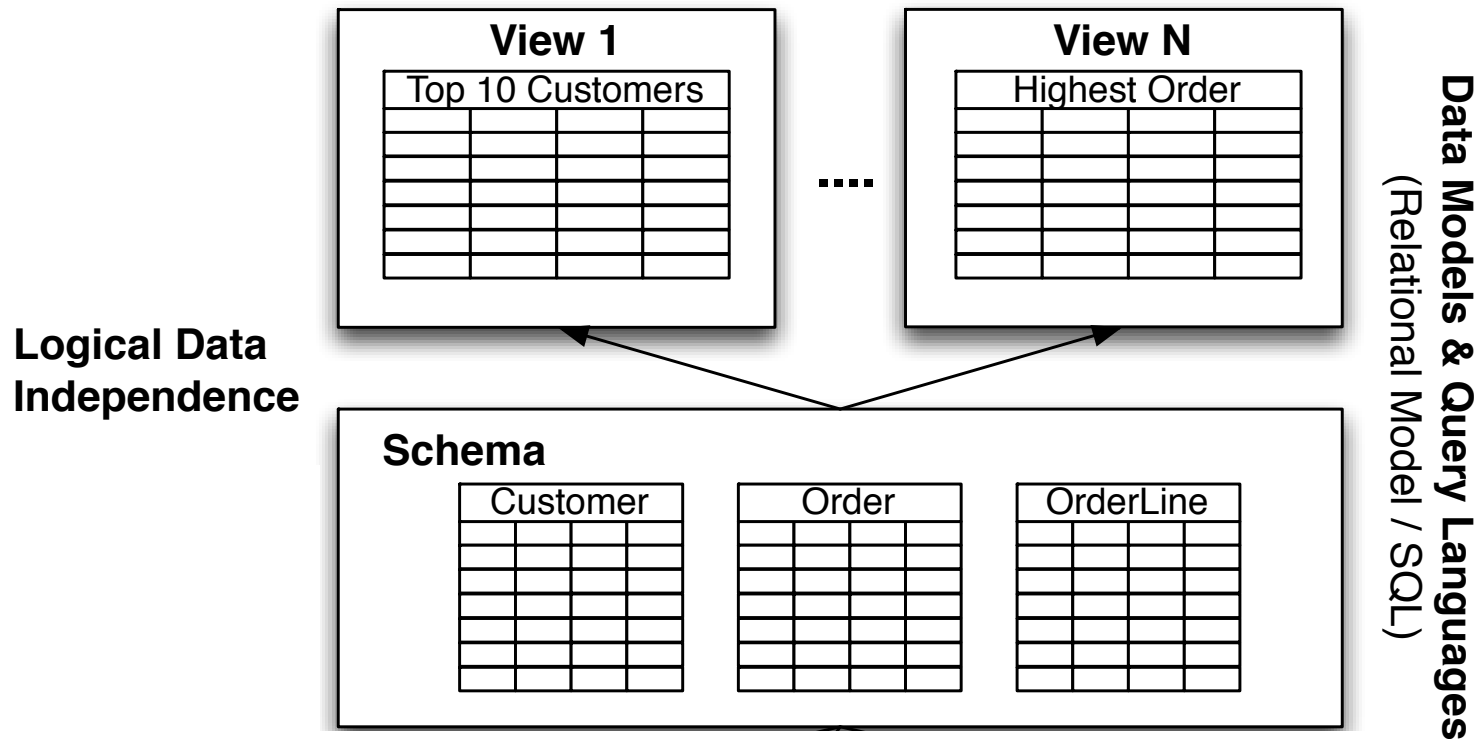
- **Scalability** → 100's of nodes
- **Data redundancy and inconsistency**
- **Data retrieval**
- **Data-Independence**

DATABASE OVERVIEW

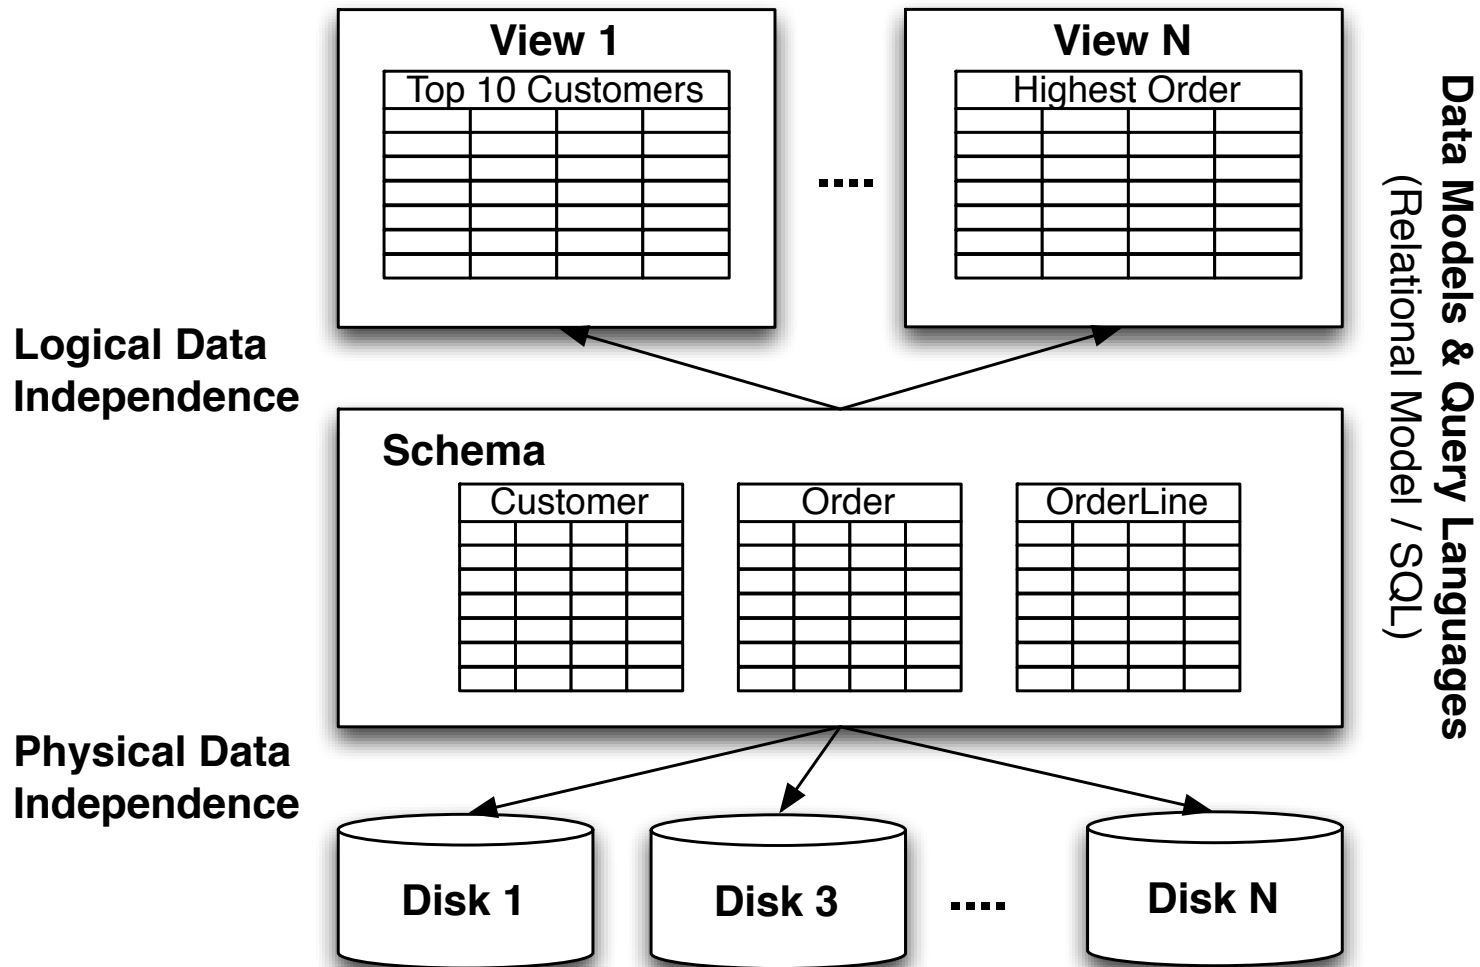
Schema

[illegible][illegible][illegible]

DATABASE OVERVIEW



DATABASE OVERVIEW

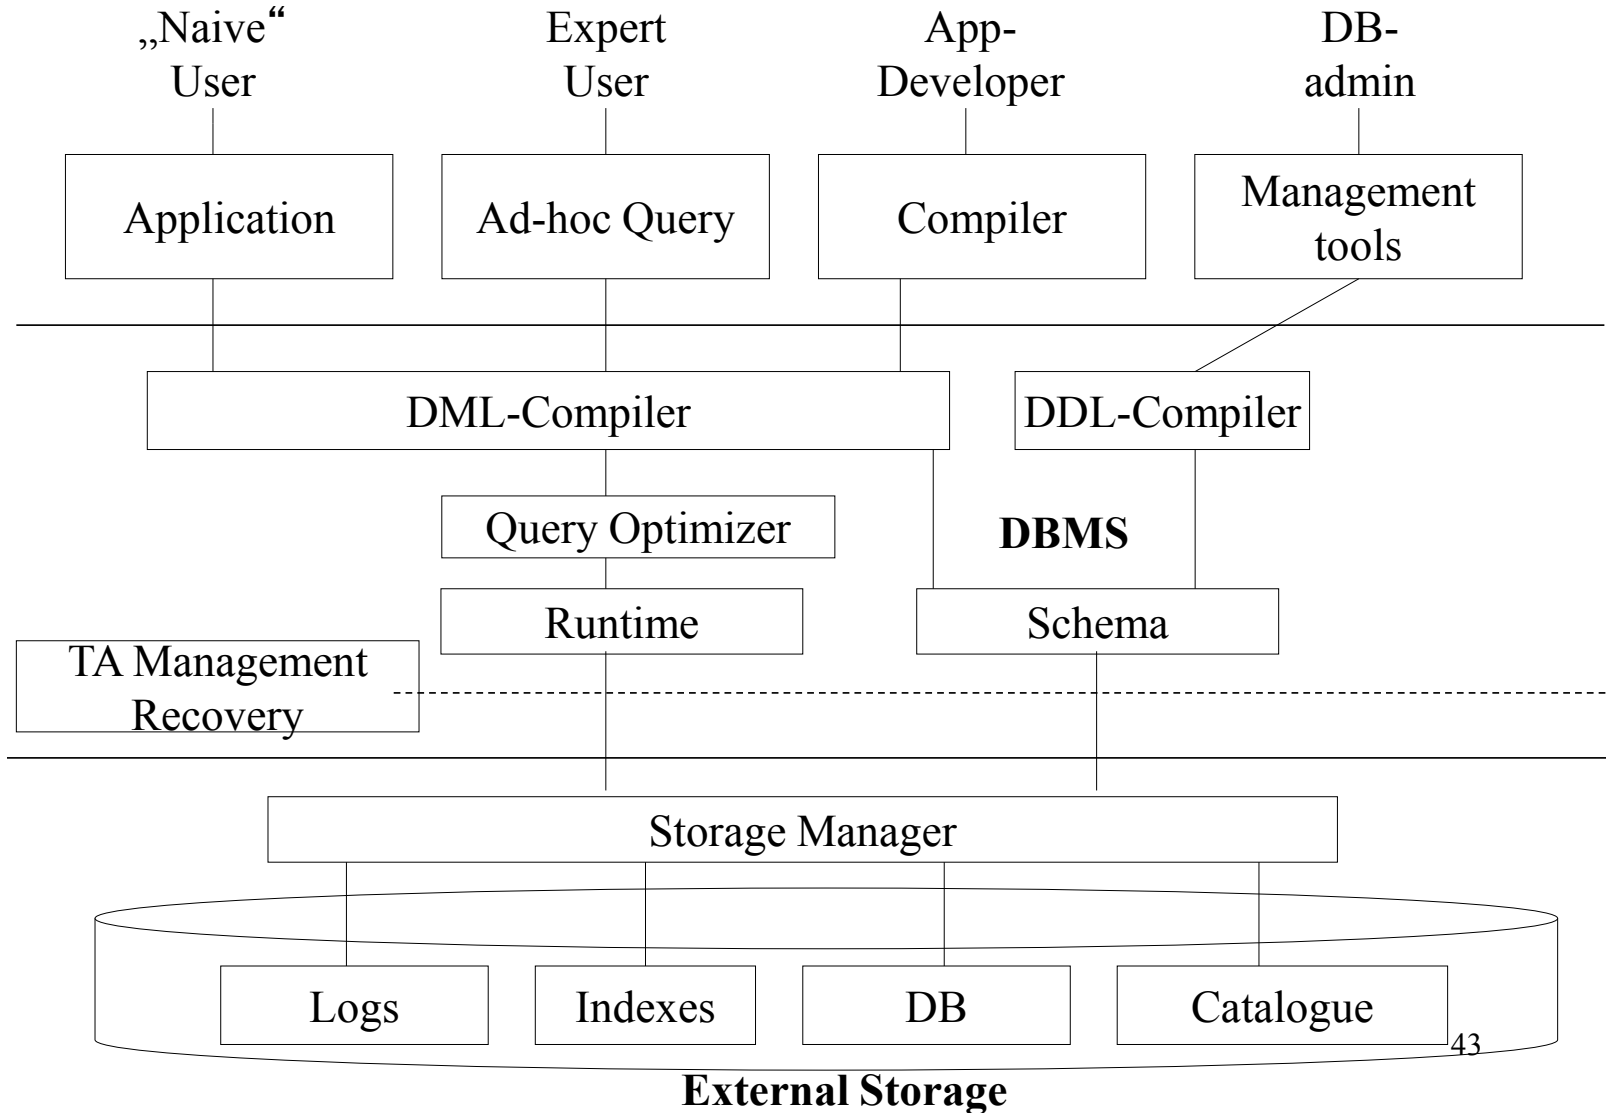


WHY DATABASES


Why not store everything in flat files?

- **Scalability** → 100's of nodes
- **Data redundancy and inconsistency**
- **Data retrieval**
- **Data-Independence**
- **Concurrent access**
- **Security and privacy**

COMPONENTS OF A DATABASE SYSTEM



DATABASES FOR DATA SCIENTIST




Book of duty

Conceptual Design
(ER)

- Logical design (schema)
- Physical design (index, hints)

DATABASES FOR DATA SCIENTIST



Book of duty

Conceptual Design
(ER)

- Logical design (schema)
- Physical design (index, hints)

BOOK OF DUTY

Describe information requirements

- **Objects used (e.g., student, professor, lecture)**
- **Domains of attributes of objects**
- **Identifiers, references / relationships**


Describe processes

- E.g., examination, degree, register course

Describe processing requirements

- Cardinalities: how many students?
- Distributions: skew of lecture attendance
- Workload: how often a process is carried out
- Priorities and service level agreements

DATABASES FOR DATA SCIENTIST



Book of duty

Conceptual Design
(ER)

- Logical design (schema)
- Physical design (index, hints)

ENTITY/RELATIONSHIP (ER) MODEL

Entity

Student

Relationship

attends

Attribute

Name

Key

Student-
ID

Student-
ID

Role

Attendant

ENTITY/RELATIONSHIP (ER) MODEL

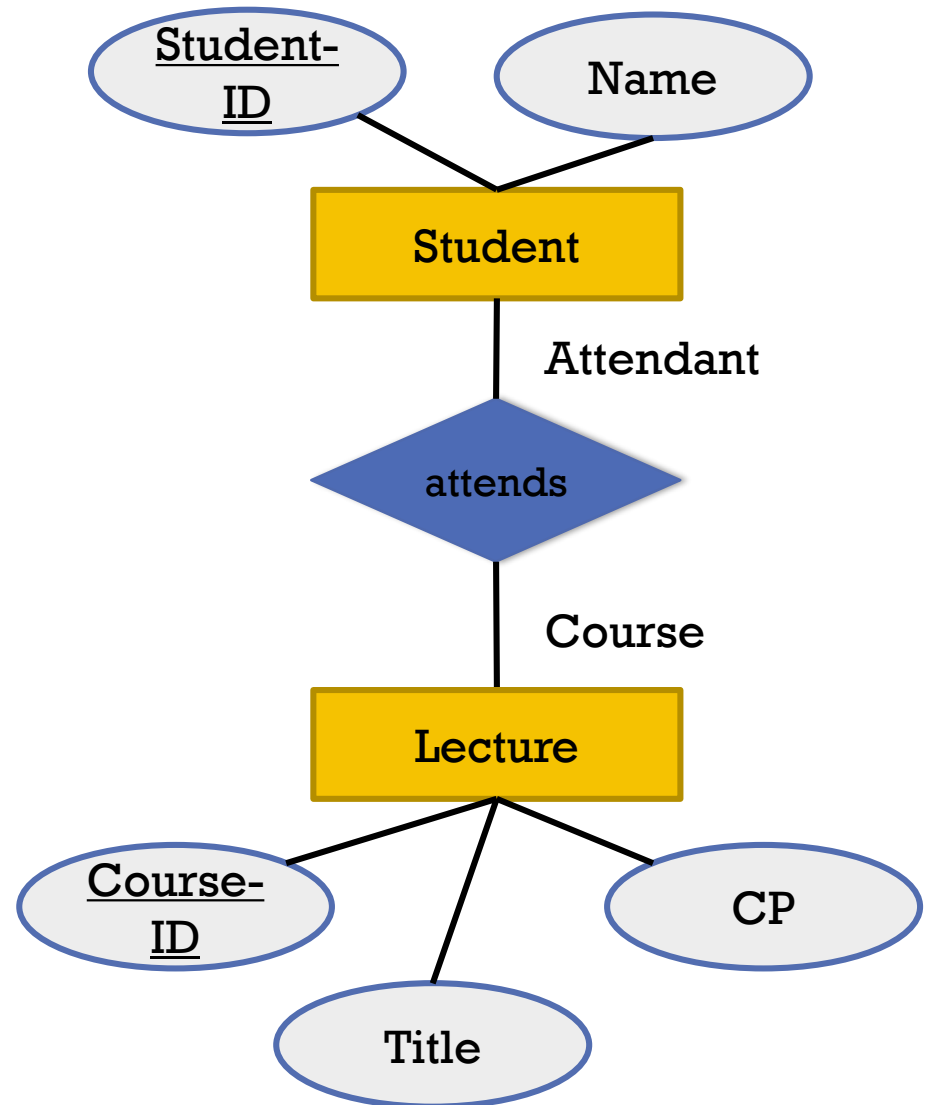
Entity

Relationship

Attribute

Key

Role



WHY ER

Advantages

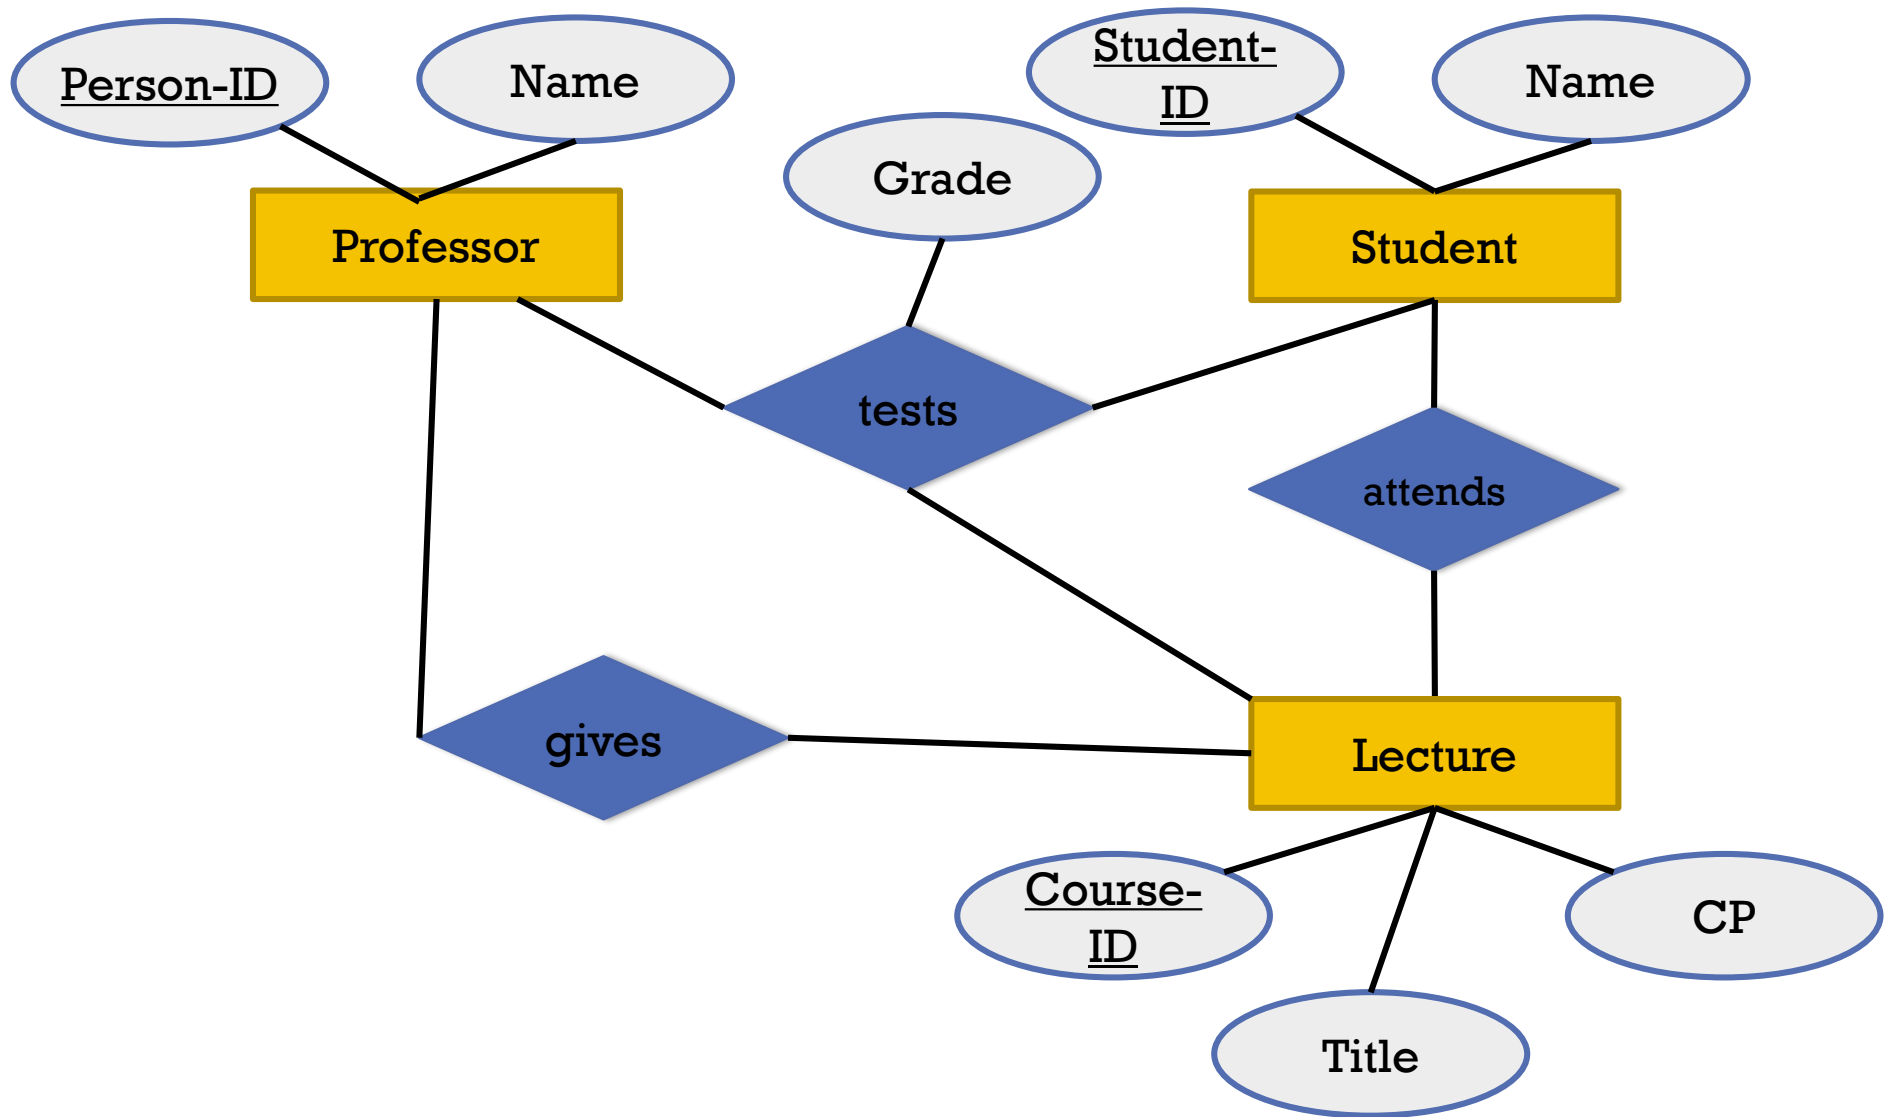
- ER diagrams are easy to create
- ER diagrams are easy to edit
- ER diagrams are easy to read (from the layman)
- ER diagrams express all information requirements

Other aspects


- Minimality
- Tools (e.g., Visio)
- Graphical representation

General





- Try to be concise, complete, comprehensible, and correct
- Controversy whether ER/UML is useful in practice
- No controversy that everybody needs to learn ER/UML



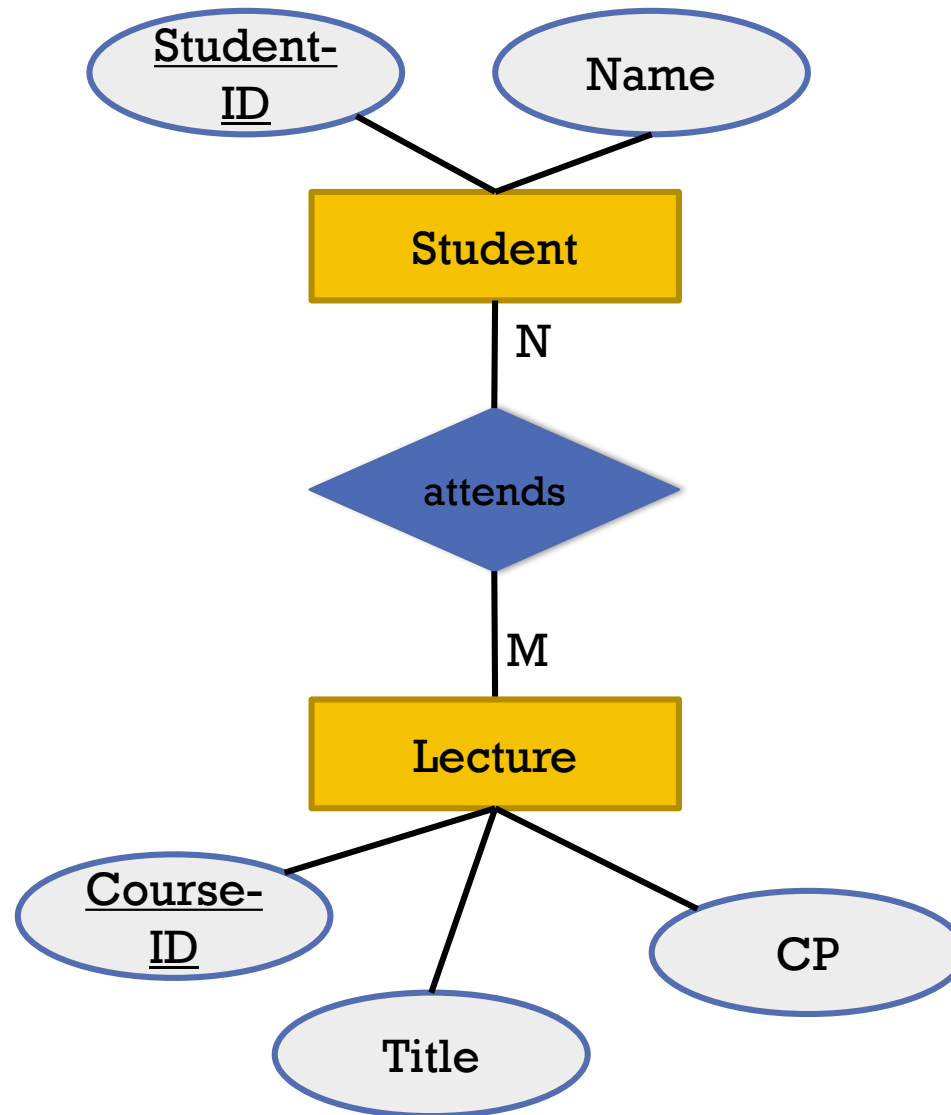
FUNCTIONALITIES



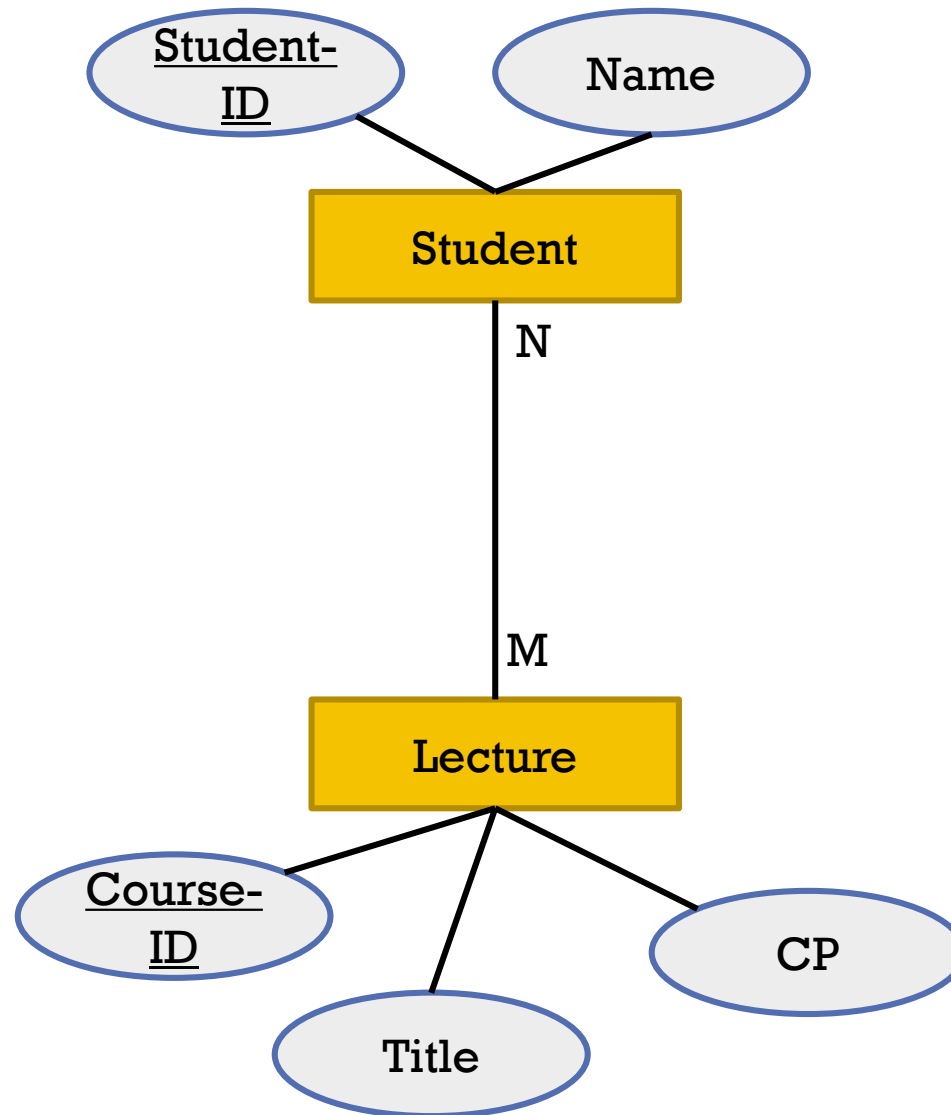
$$R \subseteq E_1 \times E_2$$




EXAMPLE: PROFESSOR <-> LECTURE



SOMETIMES ALSO SHOWN AS

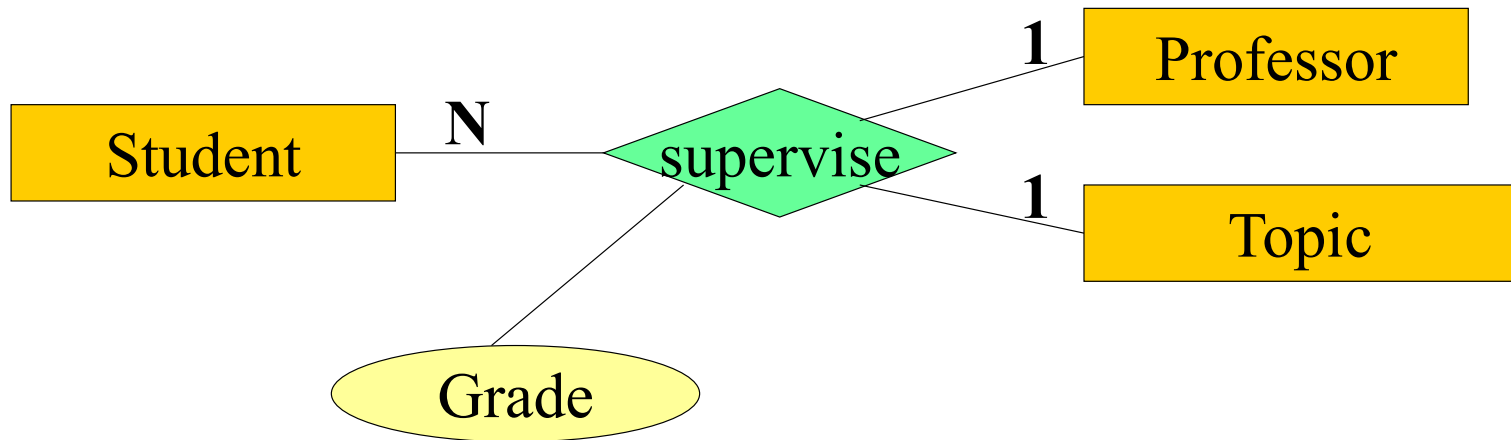


FUNCTIONALITIES OF N-ARY RELATIONSHIPS



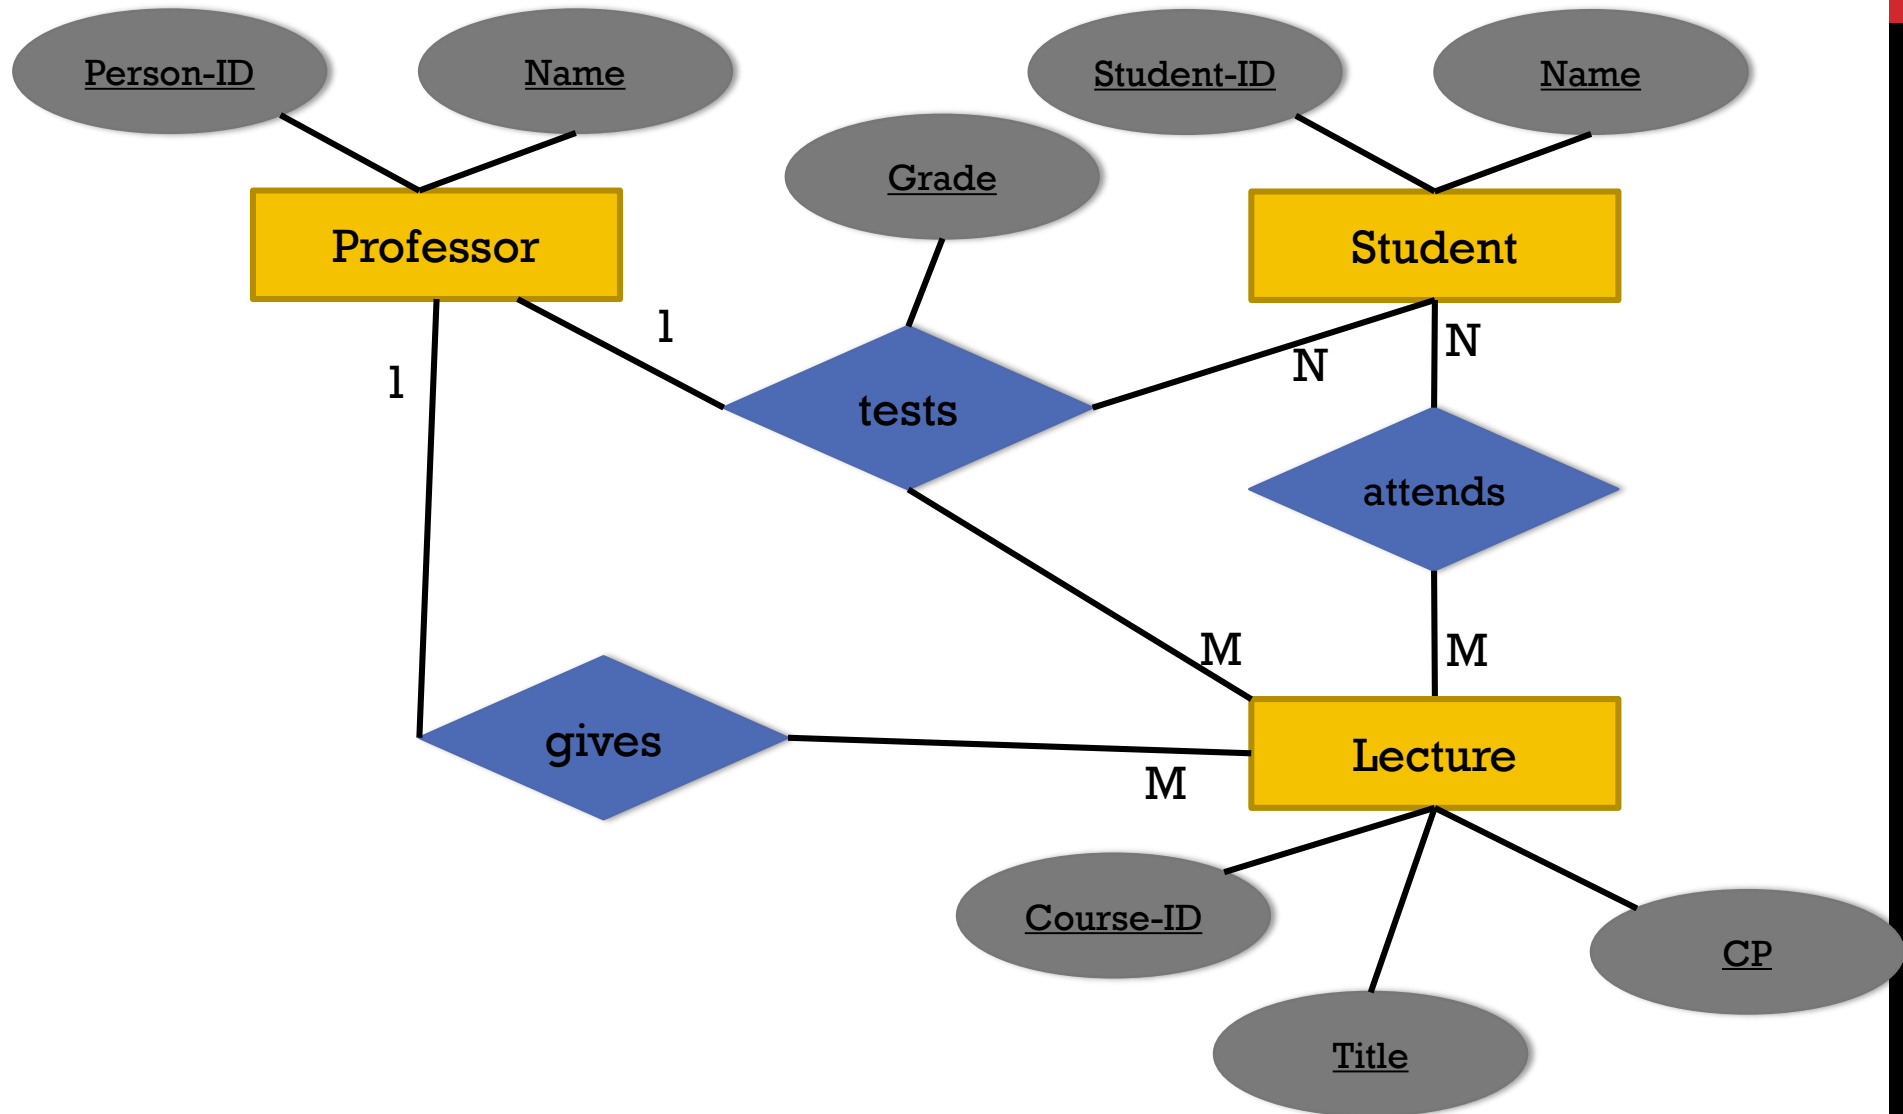
$$R : E_1 \times \dots \times E_{k-1} \times E_{k+1} \times \dots \times E_n \rightarrow E_k$$

EXAMPLE SEMINAR



$\text{supervise} : \text{Professor} \times \text{Student} \rightarrow \text{Topic}$

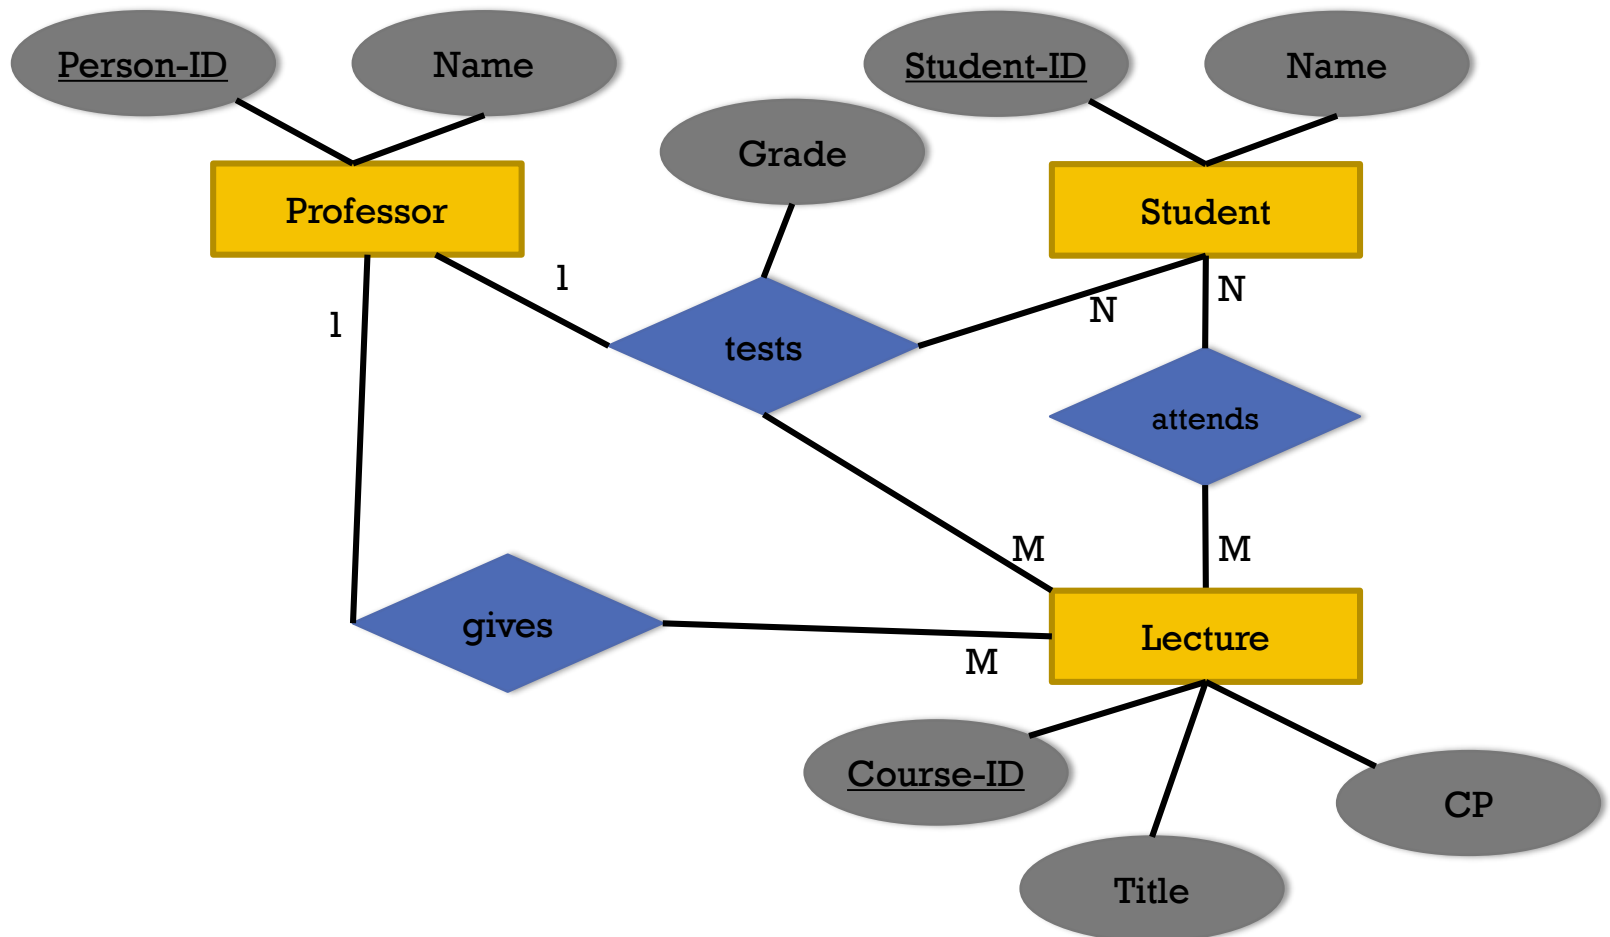
$\text{supervise} : \text{Topic} \times \text{Student} \rightarrow \text{Professor}$



ATTRIBUTE VS ENTITY

Should the *grade* be an entity or attribute?

Should test be an entity or relationship?



RULES OF THUMB

Attribute vs. Entity

- Entity if the concept has more than one relationship
- Attribute if the concept has only one 1:1 relationship

Partitioning of ER Models

- Most realistic models are larger than a page
- Partition by domains (library, research, finances, ...)

Good vs. Bad models

- Do not model redundancy or tricks to improve performance
- Less entities is better (the fewer, the better!)
- Remember the 5 C's (clear, concise, correct, complete, compliant)

NOT COVERED

(Min,Max) - Notation

Weak Entities

Generalization (i.e., inheritance)

Modeling limitations

Enhanced ERM

...

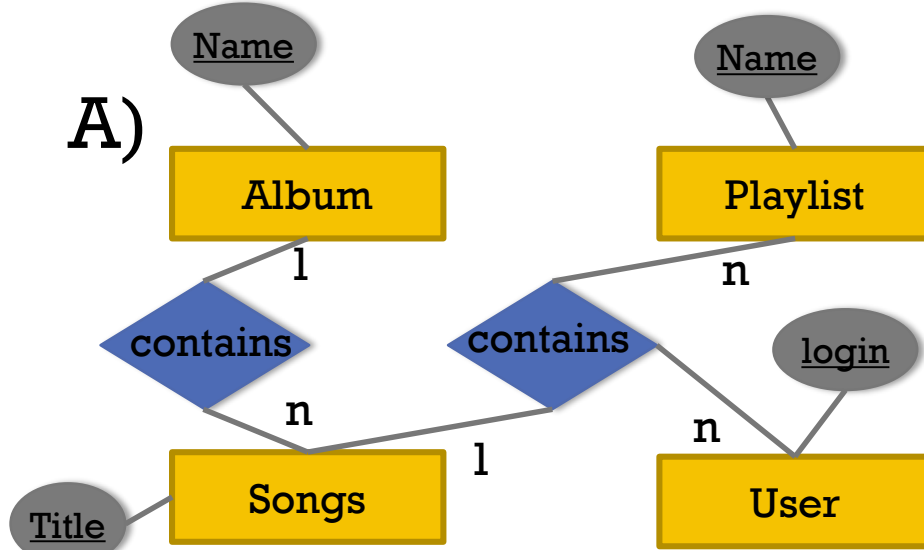
CLICKER QUESTION

Model a music record database

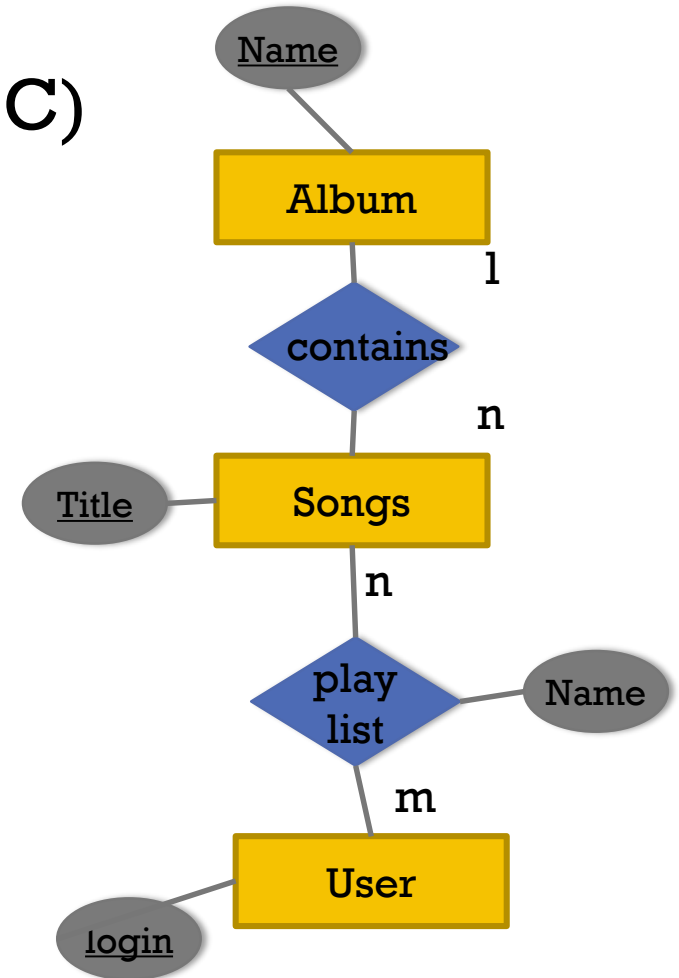
- An album has a unique name and songs have unique titles
- An album contains several songs
- A playlist has a unique name and is created by one user with a unique login
- A playlist contains several songs from potential different albums

CLICKER QUESTION

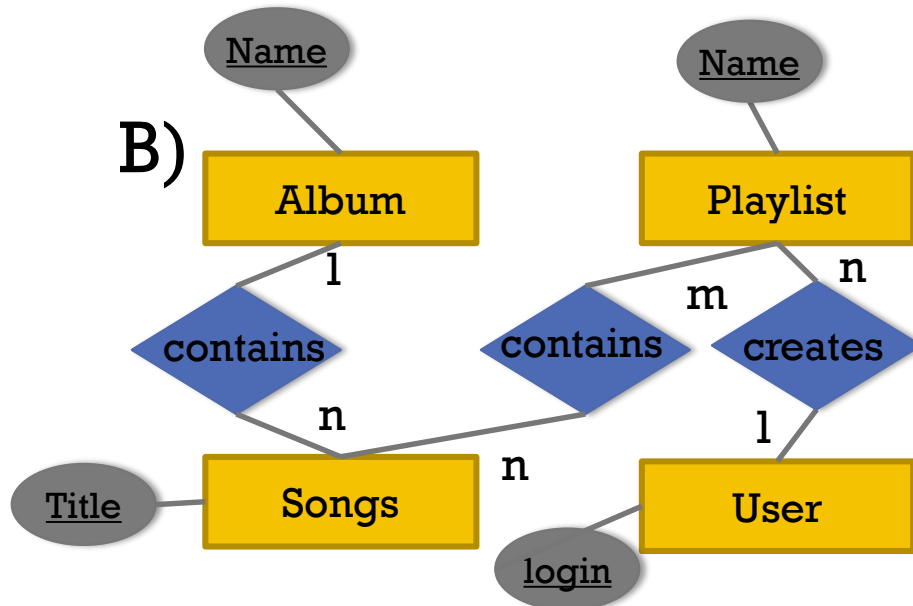
A)




C)



B)



DATABASES FOR DATA SCIENTIST



Book of duty

Conceptual Design
(ER)

- Logical design (schema)
- Physical design (index, hints)

RELATIONAL MODEL - TERMS

Account =

bname	acct_no	balance
Downtown	A-101	500
Brighton	A-201	900
Brighton	A-217	500

Table name

Attribute names

Terms

- Tables → Relations
- Columns → Attributes
- Rows → Tuples
- Schema (e.g.: Acct_Schema = (bname, acct_no, balance))

WHY ARE THEY CALLED RELATIONS?

Relation:

- $R \subseteq D_1 \times \dots \times D_n$
- D_1, D_2, \dots, D_n are domains

Example: AddressBook \subseteq string x string x integer

WHY ARE THEY CALLED RELATIONS?

Relation:

- $R \subseteq D_1 \times \dots \times D_n$
- D_1, D_2, \dots, D_n are domains

Example: AddressBook \subseteq string \times string \times integer

Tuple: $t \in R$

Example: $t = (\text{„Mickey Mouse“}, \text{„Main Street“}, 4711)$

WHY ARE THEY CALLED RELATIONS?

Relation:

- $R \subseteq D_1 \times \dots \times D_n$
- D_1, D_2, \dots, D_n are domains

Example: AddressBook \subseteq string \times string \times integer

Tuple: $t \in R$

Example: $t = (\text{„Mickey Mouse“}, \text{„Main Street“}, 4711)$

Schema: associates labels to domains

Example:

AddrBook: {[Name: string, Address: string, Tel#:integer]}

RELATIONS

bname	acct_no	balance
Downtown	A-101	500
Brighton	A-201	900
Brighton	A-217	500

Considered equivalent to...

{ (Downtown, A-101, 500),
(Brighton, A-201, 900),
(Brighton, A-217, 500) }

Relational database semantics are defined in terms of mathematical relations (i.e., sets)

KEYS AND RELATIONS

Kinds of keys

- Superkeys:
set of attributes of table for which every row has distinct set of values
- Candidate keys:
“minimal” superkeys
- Primary keys:
DBA-chosen candidate key (marked in schema by underlining)

ISBN	Title	Author	Edition	Publisher	Price
0439708184	Harry Potter	J.K. Rowling	1	Scholastic	\$6.70
0545663261	Mockingjay	<u>Suzanne</u> <u>Collins</u>	1	Scholastic	\$7.39

KEYS AND RELATIONS

Kinds of keys

- Superkeys:
set of attributes of table for which every row has distinct set of values
- Candidate keys:
“minimal” superkeys
- Primary keys:
DBA-chosen candidate key (marked in schema by underlining)

Act as Integrity Constraints

i.e., guard against illegal/invalid instance of given schema

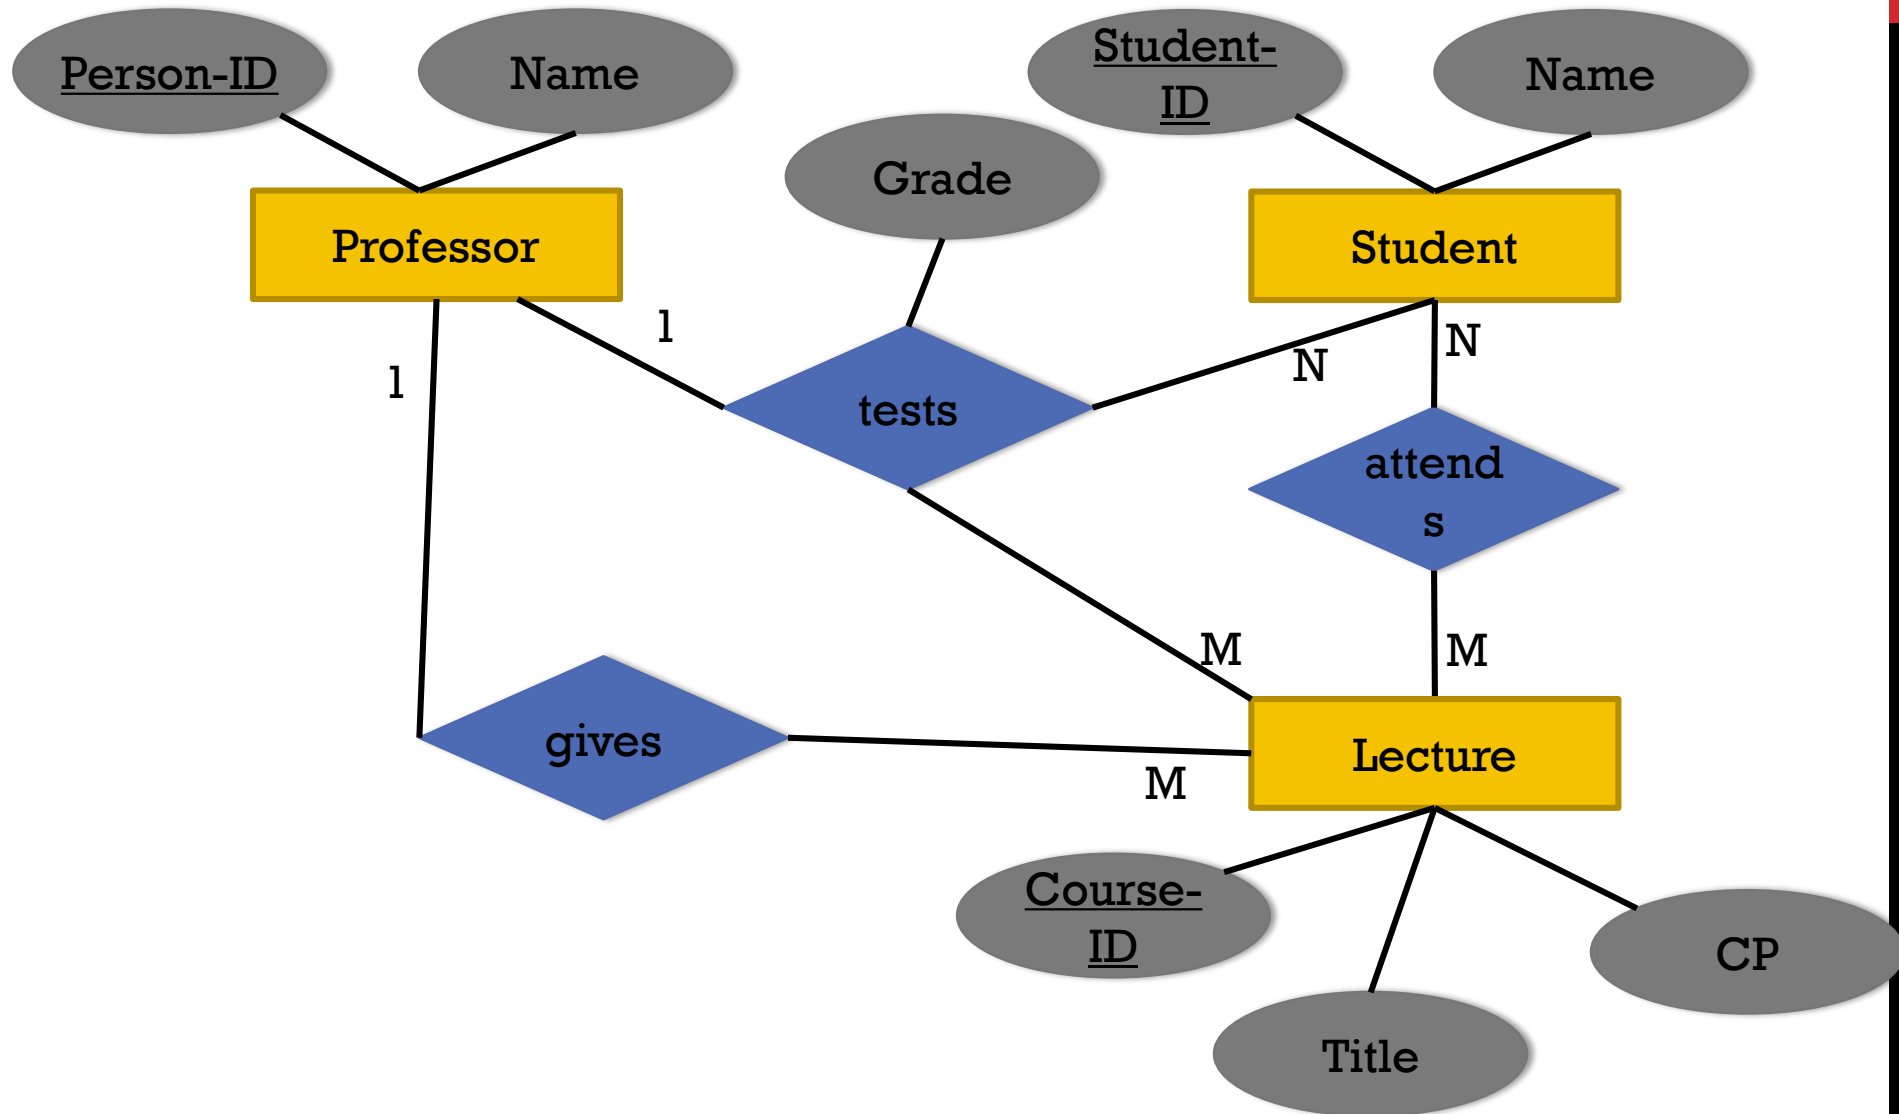
e.g., Branch = (bname, bcity, assets) \mathbb{P}

bname	bcity	assets
Brighton	Brooklyn	5M
Brighton	Boston	3M



Invalid!!

HOW TO TRANSLATE ERM TO RELATIONS

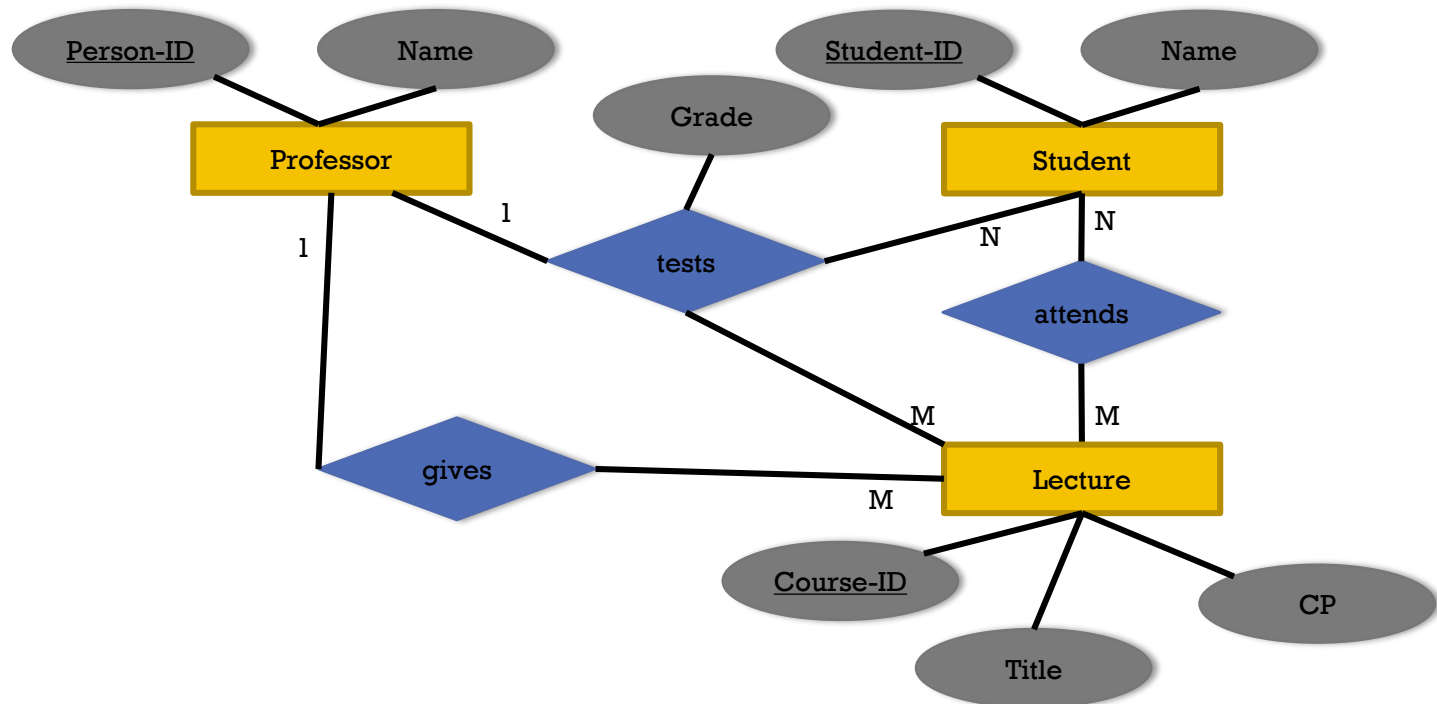


RULE #1: ENTITIES

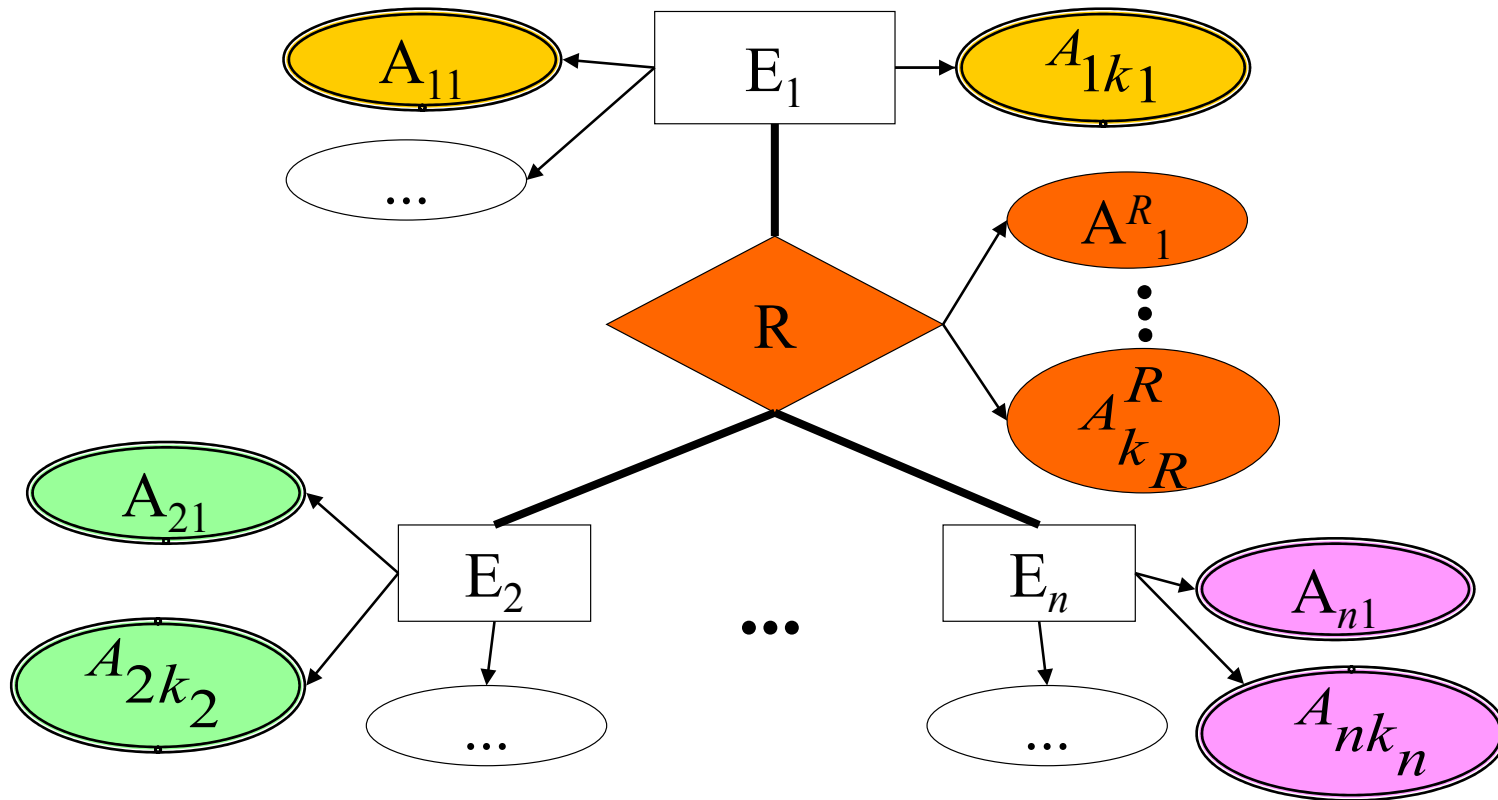
Professor(Person-ID:integer, Name:string)

Student(Student-ID:integer, Name:string)

Lecture(Course-ID:string, Title:string, CP:float)



RULE #2: RELATIONSHIPS



$$R: \left\{ \underbrace{[A_{11}, \dots, A_{1k_1}]}_{\text{Key of } E_1}, \underbrace{[A_{21}, \dots, A_{2k_2}]}_{\text{Key of } E_2}, \dots, \underbrace{[A_{n1}, \dots, A_{nk_n}]}_{\text{Key of } E_n}, \underbrace{[A_1^R, \dots, A_{k_R}^R]}_{\text{Attributes of } R} \right\}$$

RULE #2: RELATIONSHIPS

Professor(Person-ID:integer, Name:string)

Student(Student-ID:integer, Name:string)

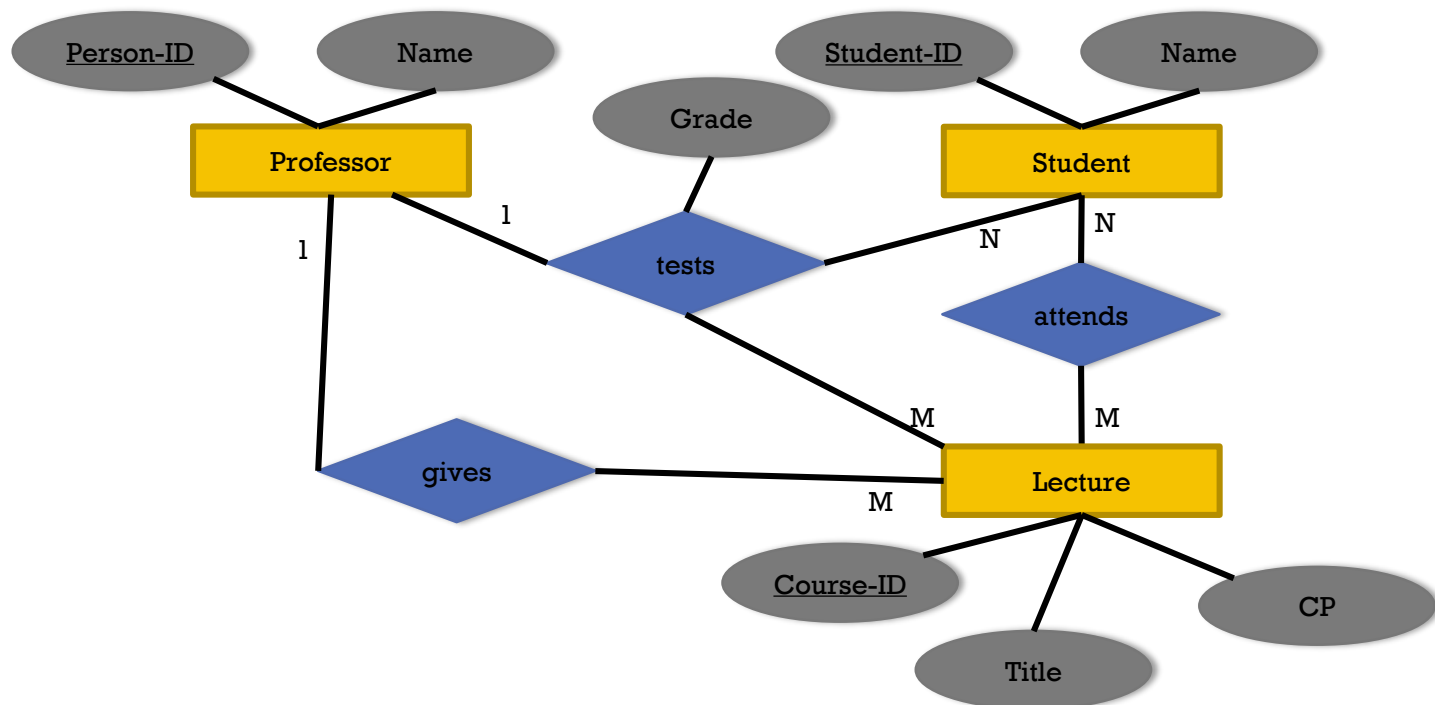
Lecture(Course-ID:string, Title:string, CP:float)

Gives(Person-ID:integer, Course-ID:string)

Attends(Student-ID:integer, Course-ID:string)

Tests(Student-ID:integer, Course-ID:string, Person-ID:integer, Grade:String)

What about keys?



RULE #2: RELATIONSHIPS

Professor(Person-ID:integer, Name:string)

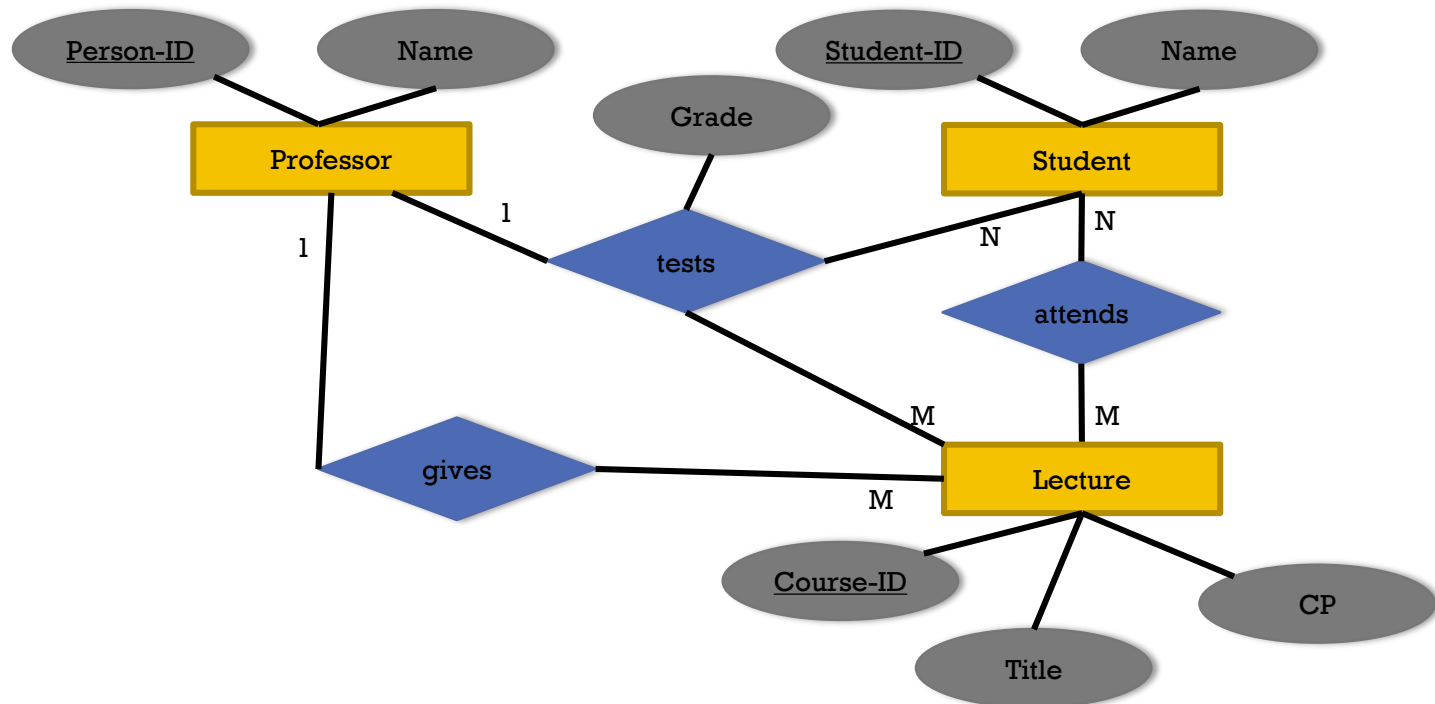
Student(Student-ID:integer, Name:string)

Lecture(Course-ID:string, Title:string, CP:float)

Gives(Person-ID:integer, Course-ID:string)

Attends(Student-ID:integer, Course-ID:string)

Tests(Student-ID:integer, Course-ID:string, Person-ID:integer, ,
Grade:string)




INSTANCE OF ATTENDS

Student	
Student-ID	...
1	...
2	...
4	...
5	...
6	...
10	...


Attends	
Student-ID	Course-ID
1	CS1951a
1	CS167
2	CS1951a
2	CS167
3	CS18
...	...

Lecture	
Course-ID	...
CS1951a	...
CS195w	...
CS18	...
CS17	...
CS142	...
CS167	...



RULE #3: MERGE RELATIONS WITH THE SAME KEY

Professor(Person-ID:integer, Name:string)
Lecture(Course-ID:string, Title:string, CP:float)
Gives(Person-ID:integer, Course-ID:string)



Professor(Person-ID:integer, Name:string)
Lecture(Course-ID:string, Title:string, CP:float, Person-ID:integer)

FINAL

```
Professor(Person-ID:integer, Name:string)
Student(Student-ID:integer, Name:string)
Lecture(Course-ID:string, Title:string, CP:float,
        Person-ID:integer)
Attends(Student-ID:integer, Course-ID:string)
Tests(Student-ID:integer, Course-ID:string,
      Person-ID:integer, Grade:string)
```


Why didn't we merge **Attends** and **Tests**?

SQL: CREATE TABLE

Professor(Person-ID:integer, Name:string)

Lecture(Course-ID:string, Title:string, CP:float, Person-ID:integer)

Student(Student-ID:integer, Name:string)



```
CREATE TABLE Student (  
  Student-ID INT,  
  Name VARCHAR(45));
```

```
CREATE TABLE Professor (  
  Person-ID INT,  
  Name VARCHAR(45));
```


```
CREATE TABLE Lecture(  
  Course-ID INT,  
  Title VARCHAR(45),  
  CP REAL,  
  Person-ID INT);
```

SQL: CREATE TABLE

Professor(Person-ID:integer, Name:string)

Lecture(Course-ID:string, Title:string, CP:float, Person-ID:integer)

Student(Student-ID:integer, Name:string)



```
CREATE TABLE Student (  
  Student-ID INT,  
  Name VARCHAR(45));
```

```
CREATE TABLE Professor (  
  Person-ID INT,  
  Name VARCHAR(45));
```

```
CREATE TABLE Lecture(  
  Course-ID INT,  
  Title VARCHAR(45),  
  CP REAL,  
  Person-ID INT);
```

INTEGRITY CONSTRAINTS IN CREATE TABLE

not null

primary key (A_1, \dots, A_n)

Example: Declare *ID* as the primary key for *instructor*

```
CREATE TABLE Attends (  
    Student-ID INT,  
    Course-ID VARCHAR(6),  
    PRIMARY KEY (Student-ID, Course-ID));
```

primary key declaration on an attribute automatically ensures **not null**

SINGLE ATTRIBUTE KEY

```
CREATE TABLE course (  
    course_id VARCHAR(8) PRIMARY KEY,  
    title VARCHAR(50),  
    cp NUMERIC(1,1));
```

Primary key declaration can be combined with attribute declaration as shown above

FOREIGN KEYS

```
CREATE TABLE `Attends` (  
  `Student_ID` INT NOT NULL,  
  `Course-ID` VARCHAR(6) NOT NULL,  
  PRIMARY KEY (`Student_ID`, `Course-ID`),  
  CONSTRAINT `fk_attend_Student`  
    FOREIGN KEY (`Student_ID`)  
    REFERENCES `Student` (`ID`)  
    ON DELETE NO ACTION  
    ON UPDATE NO ACTION,  
  CONSTRAINT `fk_attend_lecture`  
    FOREIGN KEY (`Lecture_Course-ID`)  
    REFERENCES `Lecture` (`Course-ID`)  
    ON DELETE NO ACTION  
    ON UPDATE NO ACTION;
```

INDEX

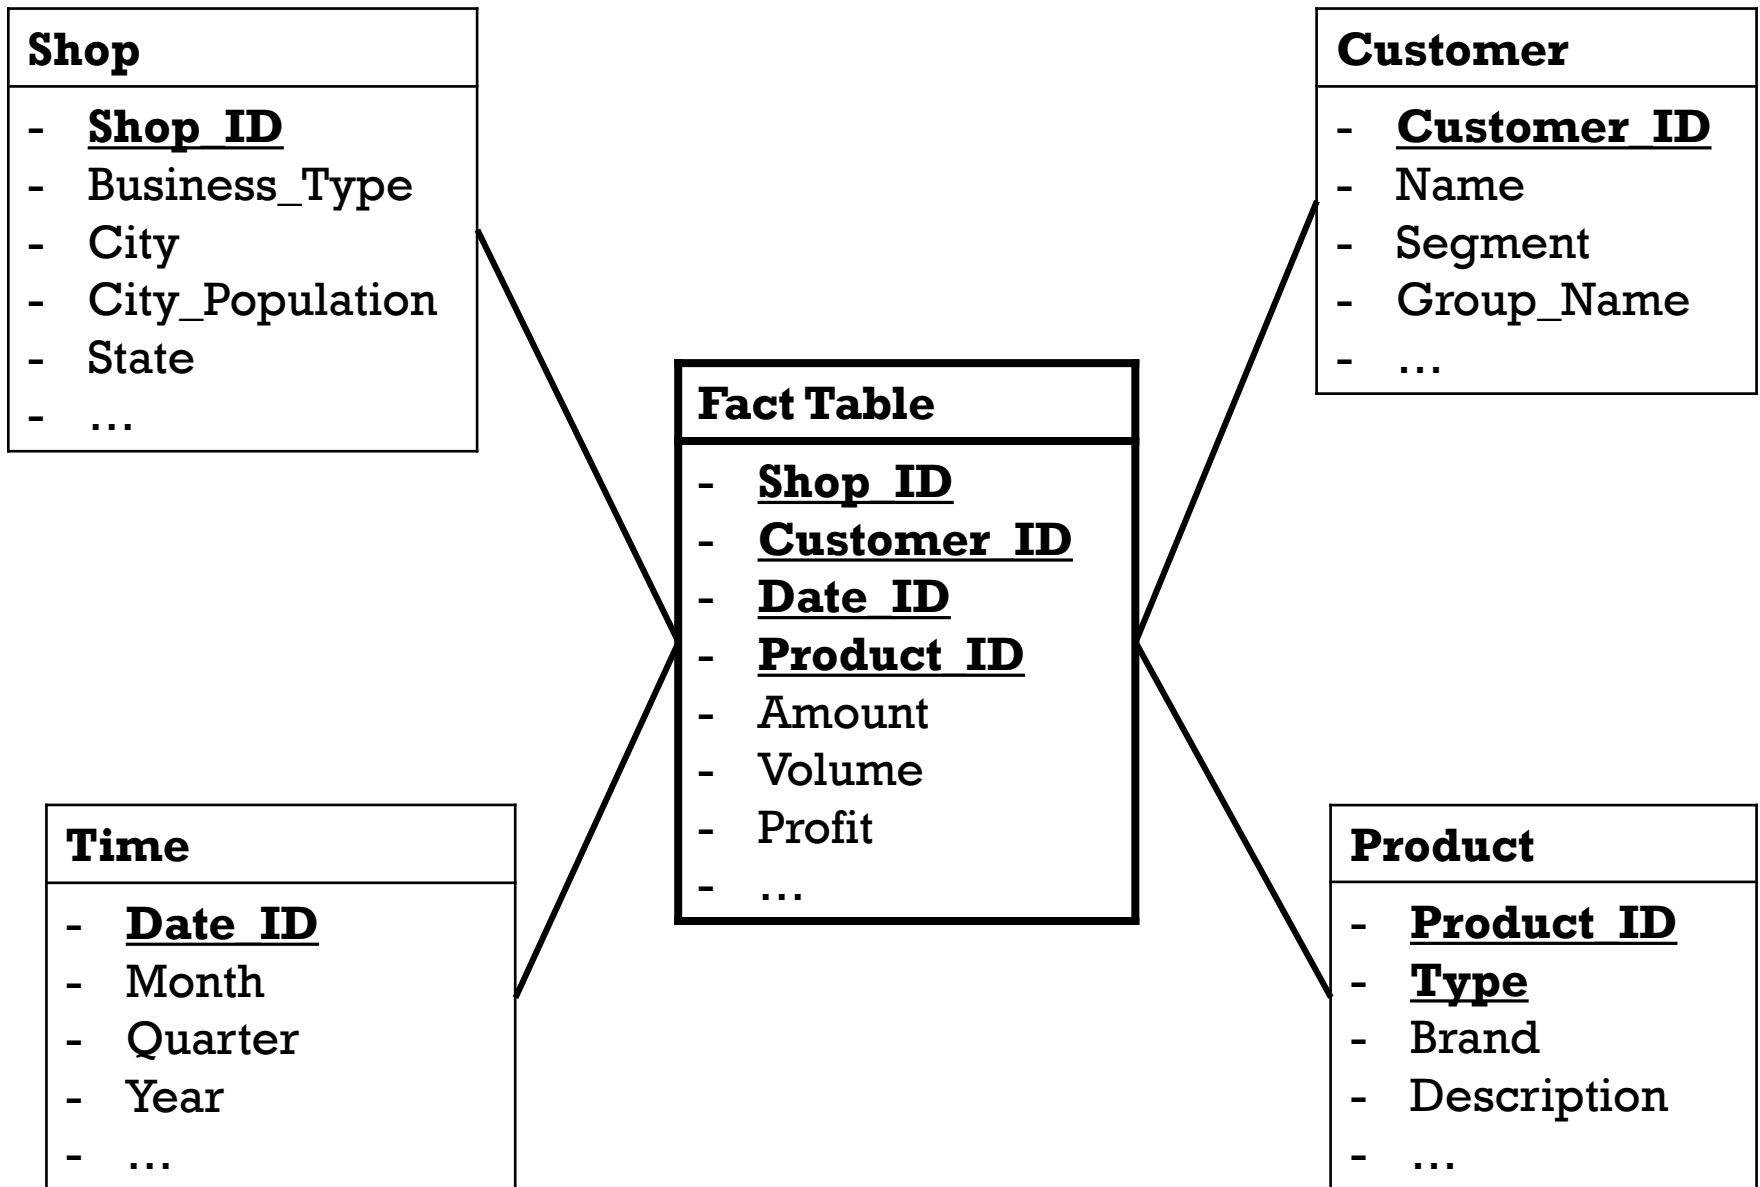
```
CREATE INDEX `fk_Student_has_Lecture_Lecture1_idx` ON `Attends`  
(`Course-ID` ASC);
```

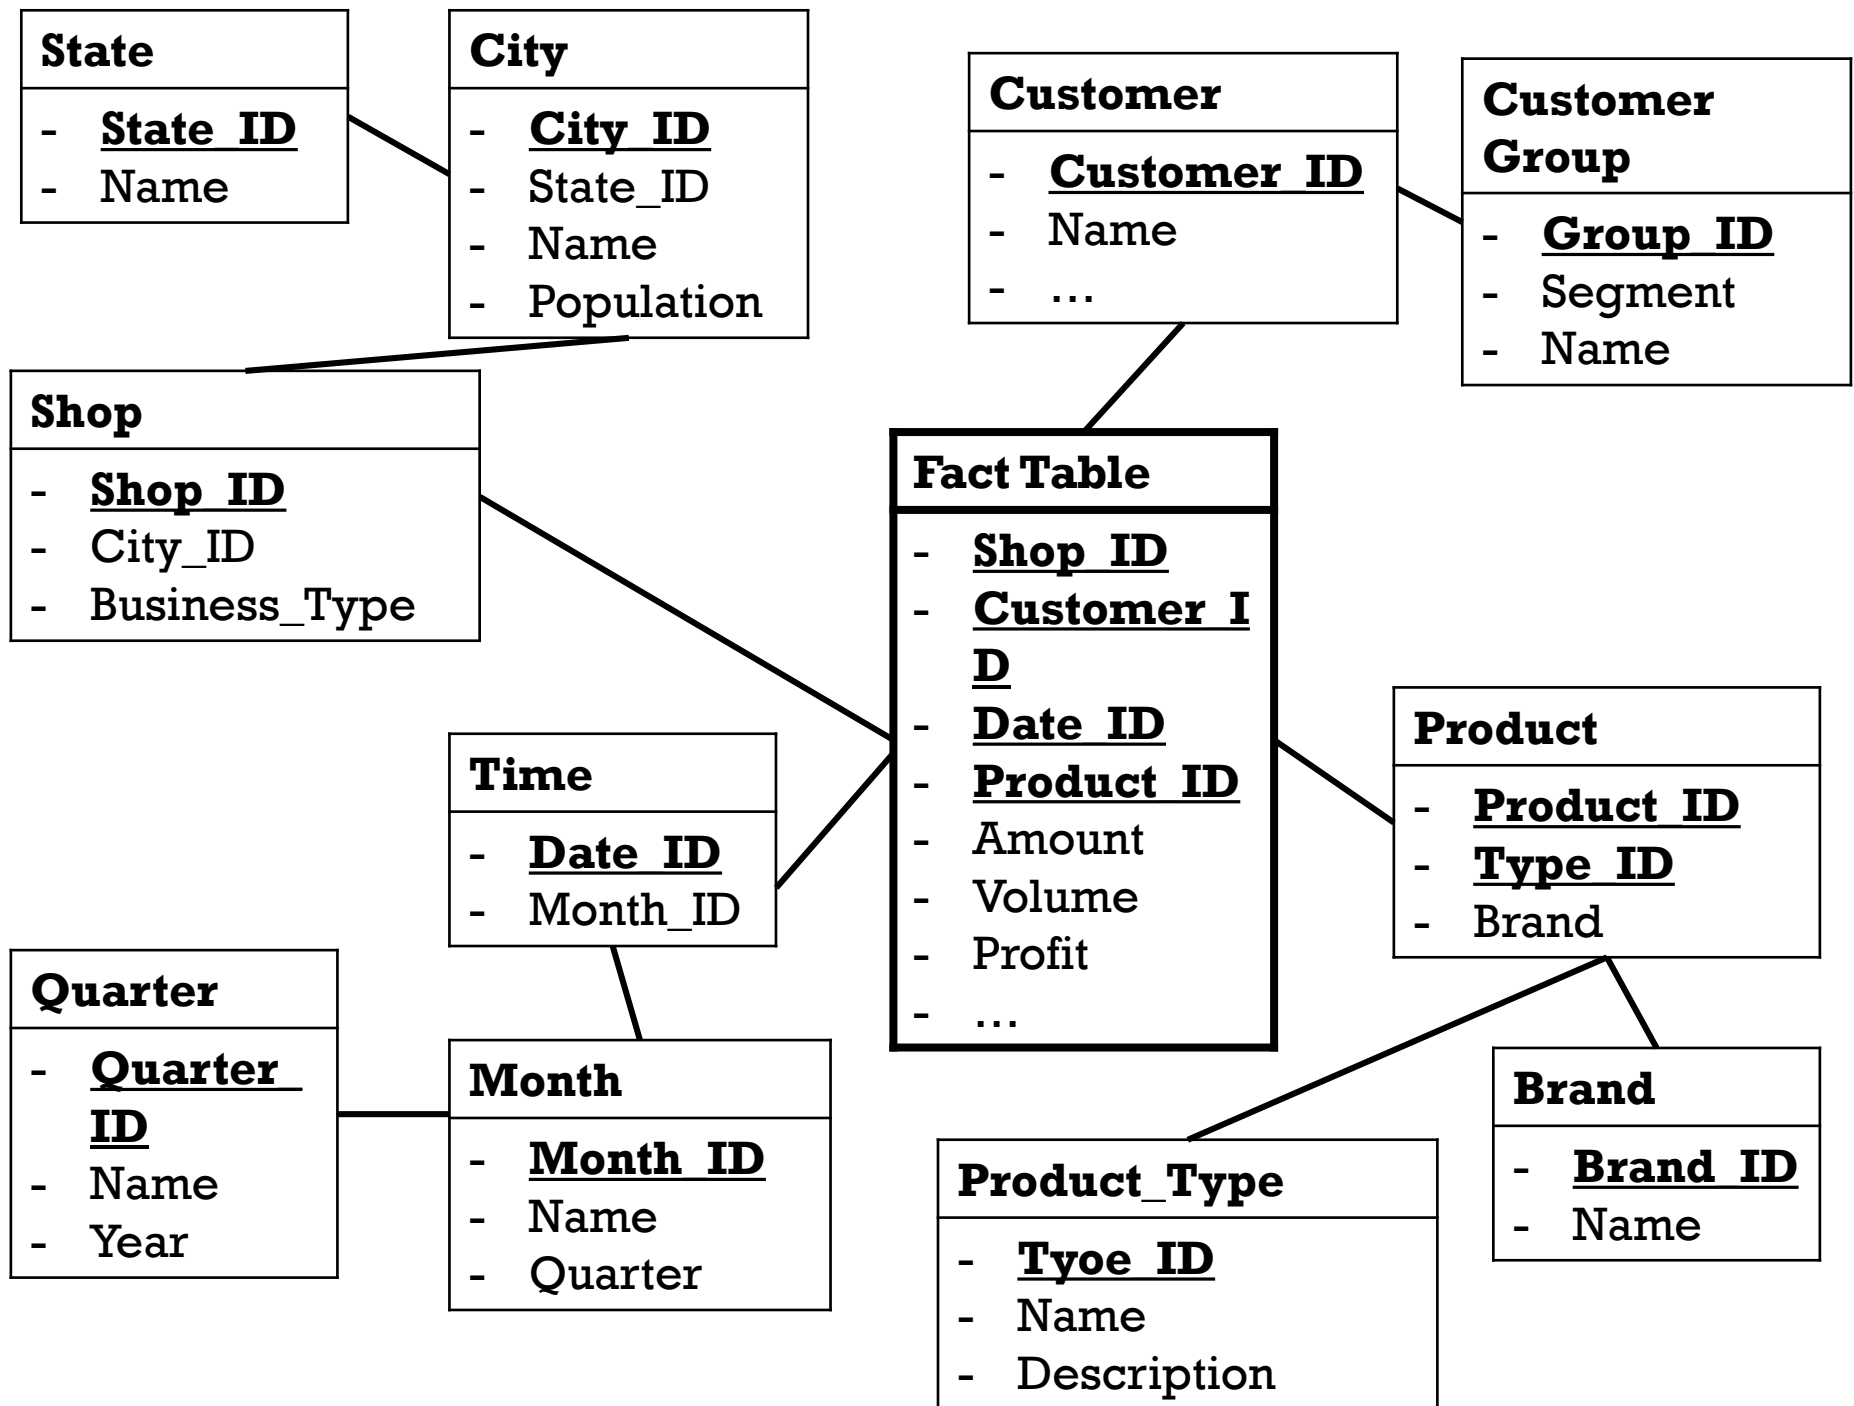
```
CREATE INDEX `fk_Student_has_Lecture_Student_idx` ON `Attends`  
(`Student-ID` ASC);
```

PROBLEM

- **You are the new Data Scientist at Evil Market**
- Evil Market is tracking all customer purchases with their membership card or credit card
- They also have data about their customers (estimated income, family status,...)
- Recently, they are trying to improve their image for young mothers
- As a start they want to know the following information for mothers under 30 for 2013:
 - How much do they spend?
 - How much do they spend per state?
 - How does this compare to all customers under 30?
 - What are their favorite products?
 - How much do they spend per year?

Your first project: Design the schema for Evil Market!






OTHER ANNOUNCEMENTS

Want to get involved in research?

We are offering several independent studies and summer research internship.

Sign-up available on: <http://database.cs.brown.edu/> or directly: <http://tinyurl.com/zxznf92>

Possible Topics:



Infiniband



Tupleware



Interactive Data
Exploration

Do you want to
drink from the fire
hose??? Then take
my
CS227 class

